

Matrix multiplication

GPU vs CPU (not transposed)

N	t_GPU [ms]	t_CPU [ms]
128	4.58	14.1
256	42.5	148.64
512	345.17	1104.6
1024	2864.53	39664.86

Workgroups 1x1

GPU vs CPU (transposed)

N	t_GPU [ms]	t_CPU [ms]
512	255.58	283.17
1024	2018.56	2522.37

Influence of Workgroup size

N_workgr	t_GPU [ms]
1	345.17
2	307.57
4	151.47
8	111.84
16	636.91

not transposed, N = 512

Exploiting cache locality is crucial!

